



In [1]:

```
import keras
import tensorflow as tf
print('TensorFlow version:', tf.__version__)
print('Keras version:', keras.__version__)
```

Using TensorFlow backend.

TensorFlow version: 1.11.0
Keras version: 2.2.4



In [2]:

```
import os
from os.path import join
import json
import random
import itertools
import re
import datetime
import cairocffi as cairo
import editdistance
import numpy as np
from scipy import ndimage
import pylab
import matplotlib.pyplot as plt
import matplotlib.gridspec as gridspec
from keras import backend as K
from keras.layers.convolutional import Conv2D, MaxPooling2D
from keras.layers import Input, Dense, Activation
from keras.layers import Reshape, Lambda
from keras.layers.merge import add, concatenate
from keras.models import Model, load_model
from keras.layers.recurrent import GRU
from keras.optimizers import SGD
from keras.utils.data_utils import get_file
from keras.preprocessing import image
import keras.callbacks
import cv2
```



In [3]:

```
sess = tf.Session()
K.set_session(sess)
```

Get alphabet



In [4]:

```

from collections import Counter
def get_counter(dirpath, tag):
    dirname = os.path.basename(dirpath)
    ann_dirpath = join(dirpath, 'ann')
    letters = ''
    lens = []
    for filename in os.listdir(ann_dirpath):
        json_filepath = join(ann_dirpath, filename)
        ann = json.load(open(json_filepath, 'r'))
        tags = ann['tags']
        if tag in tags:
            description = ann['description']
            lens.append(len(description))
            letters += description
    print('Max plate length in "%s":' % dirname, max(Counter(lens).keys()))
    return Counter(letters)
c_val = get_counter('/data/anpr_ocr__train', 'val')
c_train = get_counter('/data/anpr_ocr__train', 'train')
letters_train = set(c_train.keys())
letters_val = set(c_val.keys())
if letters_train == letters_val:
    print('Letters in train and val do match')
else:
    raise Exception()
# print(len(letters_train), len(letters_val), len(letters_val | letters_train))
letters = sorted(list(letters_train))
print('Letters:', ' '.join(letters))

```

```

Max plate length in "anpr_ocr__train": 8
Max plate length in "anpr_ocr__train": 8
Letters in train and val do match
Letters: 0 1 2 3 4 5 6 7 8 9 A B C E H K M O P T X Y

```

Input data generator



In [5]:

```

def labels_to_text(labels):
    return ''.join(list(map(lambda x: letters[int(x)], labels)))

def text_to_labels(text):
    return list(map(lambda x: letters.index(x), text))

def is_valid_str(s):
    for ch in s:
        if not ch in letters:
            return False
    return True

class TextImageGenerator:

    def __init__(self,
                  dirpath,
                  tag,
                  img_w, img_h,
                  batch_size,
                  downsample_factor,
                  max_text_len=8):

        self.img_h = img_h
        self.img_w = img_w
        self.batch_size = batch_size
        self.max_text_len = max_text_len
        self.downsample_factor = downsample_factor

        img_dirpath = join(dirpath, 'img')
        ann_dirpath = join(dirpath, 'ann')
        self.samples = []
        for filename in os.listdir(img_dirpath):
            name, ext = os.path.splitext(filename)
            if ext in ['.png', '.jpg']:
                img_filepath = join(img_dirpath, filename)
                json_filepath = join(ann_dirpath, name + '.json')
                ann = json.load(open(json_filepath, 'r'))
                description = ann['description']
                tags = ann['tags']
                if tag not in tags:
                    continue
                if is_valid_str(description):
                    self.samples.append([img_filepath, description])

        self.n = len(self.samples)
        self.indexes = list(range(self.n))
        self.cur_index = 0

    def build_data(self):
        self.imgs = np.zeros((self.n, self.img_h, self.img_w))
        self.texts = []
        for i, (img_filepath, text) in enumerate(self.samples):
            img = cv2.imread(img_filepath)
            img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
            img = cv2.resize(img, (self.img_w, self.img_h))
            img = img.astype(np.float32)
            img /= 255

```

```

        # width and height are backwards from typical Keras convention
        # because width is the time dimension when it gets fed into the RNN
        self.imgs[i, :, :] = img
        self.texts.append(text)

def get_output_size(self):
    return len(letters) + 1

def next_sample(self):
    self.cur_index += 1
    if self.cur_index >= self.n:
        self.cur_index = 0
        random.shuffle(self.indexes)
    return self.imgs[self.indexes[self.cur_index]], self.texts[self.indexes[self.cur_index]]

def next_batch(self):
    while True:
        # width and height are backwards from typical Keras convention
        # because width is the time dimension when it gets fed into the RNN
        if K.image_data_format() == 'channels_first':
            X_data = np.ones([self.batch_size, 1, self.img_w, self.img_h])
        else:
            X_data = np.ones([self.batch_size, self.img_w, self.img_h, 1])
        Y_data = np.ones([self.batch_size, self.max_text_len])
        input_length = np.ones((self.batch_size, 1)) * (self.img_w // self.down)
        label_length = np.zeros((self.batch_size, 1))
        source_str = []

        for i in range(self.batch_size):
            img, text = self.next_sample()
            img = img.T
            if K.image_data_format() == 'channels_first':
                img = np.expand_dims(img, 0)
            else:
                img = np.expand_dims(img, -1)
            X_data[i] = img
            Y_data[i] = text_to_labels(text)
            source_str.append(text)
            label_length[i] = len(text)

        inputs = {
            'the_input': X_data,
            'the_labels': Y_data,
            'input_length': input_length,
            'label_length': label_length,
            #'source_str': source_str
        }
        outputs = {'ctc': np.zeros([self.batch_size])}
        yield (inputs, outputs)

```

»

In [6]:

```

tiger = TextImageGenerator('/data/anpr_ocr__train', 'val', 128, 64, 8, 4)
tiger.build_data()

```



In [7]:

```


for inp, out in tiger.next_batch():
    print('Text generator output (data which will be fed into the neutral network):')
    print('1) the_input (image)')
    if K.image_data_format() == 'channels_first':
        img = inp['the_input'][0, 0, :, :]
    else:
        img = inp['the_input'][0, :, :, 0]

    plt.imshow(img.T, cmap='gray')
    plt.show()
    print('2) the_labels (plate number): %s is encoded as %s' %
          (labels_to_text(inp['the_labels'][0]), list(map(int, inp['the_labels'][0]
    print('3) input_length (width of image that is fed to the loss function): %d ==
          (inp['input_length'][0], tiger.img_w))
    print('4) label_length (length of plate number): %d' % inp['label_length'][0])
    break

```

Text generator output (data which will be fed into the neutral network):

1) the_input (image)



2) the_labels (plate number): P383PX65 is encoded as [18, 3, 8, 3, 18, 20, 6, 5]

3) input_length (width of image that is fed to the loss function): 30 == 128 / 4 - 2

4) label_length (length of plate number): 8

Loss and train functions, network architecture



In [8]:

```
def ctc_lambda_func(args):
    y_pred, labels, input_length, label_length = args
    # the 2 is critical here since the first couple outputs of the RNN
    # tend to be garbage:
    y_pred = y_pred[:, 2:, :]
    return K.ctc_batch_cost(labels, y_pred, input_length, label_length)

def train(img_w, load=False):
    # Input Parameters
    img_h = 64

    # Network parameters
    conv_filters = 16
    kernel_size = (3, 3)
    pool_size = 2
    time_dense_size = 32
    rnn_size = 512

    if K.image_data_format() == 'channels_first':
        input_shape = (1, img_w, img_h)
    else:
        input_shape = (img_w, img_h, 1)

    batch_size = 32
    downsample_factor = pool_size ** 2
    tiger_train = TextImageGenerator('/data/anpr_ocr__train', 'train', img_w, img_h)
    tiger_train.build_data()
    tiger_val = TextImageGenerator('/data/anpr_ocr__train', 'val', img_w, img_h, batch_size)
    tiger_val.build_data()

    act = 'relu'
    input_data = Input(name='the_input', shape=input_shape, dtype='float32')
    inner = Conv2D(conv_filters, kernel_size, padding='same',
                  activation=act, kernel_initializer='he_normal',
                  name='conv1')(input_data)
    inner = MaxPooling2D(pool_size=(pool_size, pool_size), name='max1')(inner)
    inner = Conv2D(conv_filters, kernel_size, padding='same',
                  activation=act, kernel_initializer='he_normal',
                  name='conv2')(inner)
    inner = MaxPooling2D(pool_size=(pool_size, pool_size), name='max2')(inner)

    conv_to_rnn_dims = (img_w // (pool_size ** 2), (img_h // (pool_size ** 2)) * conv_filters)
    inner = Reshape(target_shape=conv_to_rnn_dims, name='reshape')(inner)

    # cuts down input size going into RNN:
    inner = Dense(time_dense_size, activation=act, name='dense1')(inner)

    # Two layers of bidirectional GRUs
    # GRU seems to work as well, if not better than LSTM:
    gru_1 = GRU(rnn_size, return_sequences=True, kernel_initializer='he_normal', name='gru1',
                go_backwards=True)
    gru_1b = GRU(rnn_size, return_sequences=True, kernel_initializer='he_normal', name='gru1b',
                 go_backwards=False)
    gru1_merged = add([gru_1(inner), gru_1b(inner)])
    gru_2 = GRU(rnn_size, return_sequences=True, kernel_initializer='he_normal', name='gru2',
                go_backwards=True)
    gru_2b = GRU(rnn_size, return_sequences=True, kernel_initializer='he_normal', name='gru2b',
                 go_backwards=False)
    gru2_merged = add([gru_2(inner), gru_2b(inner)])

    # transforms RNN output to character activations:
```

```

inner = Dense(tiger_train.get_output_size(), kernel_initializer='he_normal',
              name='dense2')(concatenate([gru_2, gru_2b]))
y_pred = Activation('softmax', name='softmax')(inner)
Model(inputs=input_data, outputs=y_pred).summary()

labels = Input(name='the_labels', shape=[tiger_train.max_text_len], dtype='float32')
input_length = Input(name='input_length', shape=[1], dtype='int64')
label_length = Input(name='label_length', shape=[1], dtype='int64')
# Keras doesn't currently support loss funcs with extra parameters
# so CTC loss is implemented in a lambda layer
loss_out = Lambda(ctc_lambda_func, output_shape=(1,), name='ctc')([y_pred, labels, input_length, label_length])

# clipnorm seems to speeds up convergence
sgd = SGD(lr=0.02, decay=1e-6, momentum=0.9, nesterov=True, clipnorm=5)

if load:
    model = load_model('/data/mplate2_model.h5', compile=False)
else:
    model = Model(inputs=[input_data, labels, input_length, label_length], outputs=loss_out)

# the loss calc occurs elsewhere, so use a dummy lambda func for the loss
model.compile(loss={'ctc': lambda y_true, y_pred: y_pred}, optimizer=sgd)

if not load:
    # captures output of softmax so we can decode the output during visualization
    test_func = K.function([input_data], [y_pred])

    model.fit_generator(generator=tiger_train.next_batch(),
                       steps_per_epoch=tiger_train.n,
                       epochs=1,
                       validation_data=tiger_val.next_batch(),
                       validation_steps=tiger_val.n)

return model

```

Model description and training

Next block will take about 30 minutes.



In [9]:

```
model = train(128, load=True)
```

Layer (type) cted to	Output Shape	Param #	Conne
=====			
the_input (InputLayer)	(None, 128, 64, 1)	0	
=====			
conv1 (Conv2D) nput[0][0]	(None, 128, 64, 16)	160	the_i
=====			
max1 (MaxPooling2D) [0][0]	(None, 64, 32, 16)	0	conv1
=====			
conv2 (Conv2D) [0][0]	(None, 64, 32, 16)	2320	max1
=====			
max2 (MaxPooling2D) [0][0]	(None, 32, 16, 16)	0	conv2
=====			
reshape (Reshape) [0][0]	(None, 32, 256)	0	max2
=====			
dense1 (Dense) pe[0][0]	(None, 32, 32)	8224	resha
=====			
gru1 (GRU) 1[0][0]	(None, 32, 512)	837120	dense
=====			
gru1_b (GRU) 1[0][0]	(None, 32, 512)	837120	dense
=====			
add_1 (Add) [0][0]	(None, 32, 512)	0	gru1 gru1_
b[0][0]			
=====			
gru2 (GRU) [0][0]	(None, 32, 512)	1574400	add_1
=====			
gru2_b (GRU) [0][0]	(None, 32, 512)	1574400	add_1

concatenate_1 (Concatenate) [0][0]	(None, 32, 1024)	0	gru2
b[0][0]			gru2_
dense2 (Dense) tenate_1[0][0]	(None, 32, 23)	23575	conca
softmax (Activation) 2[0][0]	(None, 32, 23)	0	dense
=====			
=====			
Total params: 4,857,319			
Trainable params: 4,857,319			
Non-trainable params: 0			

Function to decode neural network output

►

In [10]:

```
# For a real OCR application, this should be beam search with a dictionary
# and language model. For this example, best path is sufficient.

def decode_batch(out):
    ret = []
    for j in range(out.shape[0]):
        out_best = list(np.argmax(out[j, 2:], 1))
        out_best = [k for k, g in itertools.groupby(out_best)]
        outstr = ''
        for c in out_best:
            if c < len(letters):
                outstr += letters[c]
        ret.append(outstr)
    return ret
```

Test on validation images



In [11]:

```

tiger_test = TextImageGenerator('/data/anpr_ocr_test', 'test', 128, 64, 8, 4)
tiger_test.build_data()

net_inp = model.get_layer(name='the_input').input
net_out = model.get_layer(name='softmax').output

for inp_value, _ in tiger_test.next_batch():
    bs = inp_value['the_input'].shape[0]
    X_data = inp_value['the_input']
    print ({net_inp:X_data})
    net_out_value = sess.run(net_out, feed_dict={net_inp:X_data})
    pred_texts = decode_batch(net_out_value)
    labels = inp_value['the_labels']
    texts = []
    for label in labels:
        text = ''.join(list(map(lambda x: letters[int(x)], label)))
        texts.append(text)

    for i in range(bs):
        fig = plt.figure(figsize=(10, 10))
        outer = gridspec.GridSpec(2, 1, wspace=10, hspace=0.1)
        ax1 = plt.Subplot(fig, outer[0])
        fig.add_subplot(ax1)
        ax2 = plt.Subplot(fig, outer[1])
        fig.add_subplot(ax2)
        print('Predicted: %s\nTrue: %s' % (pred_texts[i], texts[i]))
        img = X_data[i][:, :, 0].T
        ax1.set_title('Input img')
        ax1.imshow(img, cmap='gray')
        ax1.set_xticks([])
        ax1.set_yticks([])
        ax2.set_title('Activations')
        ax2.imshow(net_out_value[i].T, cmap='binary', interpolation='nearest')
        ax2.set_yticks(list(range(len(letters) + 1)))
        ax2.set_yticklabels(letters + ['blank'])
        ax2.grid(False)
        for h in np.arange(-0.5, len(letters) + 1 + 0.5, 1):
            ax2.axhline(h, linestyle='--', color='k', alpha=0.5, linewidth=1)

        #ax.axvline(x, linestyle='--', color='k')
        plt.show()
    break

```

```

{<tf.Tensor 'the_input_1:0' shape=(?, 128, 64, 1) dtype=float32>: array(
[[[[[0.79215688],
      [0.63921571],
      [0.36470589],
      ...,
      [0.57647061],
      [0.67843139],
      [0.73725492]],
      [0.47058824],
      [0.38039216],
      [0.22352941],
      ...,

```

```
[0.23529412],
[0.53725493],
[0.70588237]],

[[0.37254903],
 [0.29019609],
 [0.14509805],
 ...,
 [0.19607843],
 [0.50980395],
 [0.68235296]],

...,

[[0.37254903],
 [0.29019609],
 [0.14509805],
 ...,
 [0.19607843],
 [0.50980395],
 [0.68235296]],

[[0.47058824],
 [0.38039216],
 [0.22352941],
 ...,
 [0.23529412],
 [0.53725493],
 [0.70588237]],

[[0.79215688],
 [0.63921571],
 [0.36470589],
 ...,
 [0.57647061],
 [0.67843139],
 [0.73725492]]],

[[[0.79215688],
 [0.63921571],
 [0.36470589],
 ...,
 [0.57647061],
 [0.67843139],
 [0.73725492]],

[[0.47058824],
 [0.38039216],
 [0.22352941],
 ...,
 [0.23529412],
 [0.53725493],
 [0.70588237]],

[[0.37254903],
 [0.29019609],
 [0.14509805],
 ...,
 [0.19607843],
 [0.50980395],
```

```
[0.68235296]],  
  
... ,  
  
[[0.37254903],  
 [0.29019609],  
 [0.14509805],  
 ... ,  
 [0.19607843],  
 [0.50980395],  
 [0.68235296]],  
  
[[0.47058824],  
 [0.38039216],  
 [0.22352941],  
 ... ,  
 [0.23529412],  
 [0.53725493],  
 [0.70588237]],  
  
[[0.79215688],  
 [0.63921571],  
 [0.36470589],  
 ... ,  
 [0.57647061],  
 [0.67843139],  
 [0.73725492]]],  
  
[[[0.79215688],  
 [0.63921571],  
 [0.36470589],  
 ... ,  
 [0.57647061],  
 [0.67843139],  
 [0.73725492]],  
  
[[0.47058824],  
 [0.38039216],  
 [0.22352941],  
 ... ,  
 [0.23529412],  
 [0.53725493],  
 [0.70588237]],  
  
[[0.37254903],  
 [0.29019609],  
 [0.14509805],  
 ... ,  
 [0.19607843],  
 [0.50980395],  
 [0.68235296]],  
  
... ,  
  
[[0.37254903],  
 [0.29019609],  
 [0.14509805],  
 ... ,  
 [0.19607843],  
 [0.50980395],
```

```

[0.68235296]],

[[0.47058824],
 [0.38039216],
 [0.22352941],
 ...,
 [0.23529412],
 [0.53725493],
 [0.70588237]],

[[0.79215688],
 [0.63921571],
 [0.36470589],
 ...,
 [0.57647061],
 [0.67843139],
 [0.73725492]]],

...,

[[[0.79215688],
   [0.63921571],
   [0.36470589],
   ...,
   [0.57647061],
   [0.67843139],
   [0.73725492]],

 [0.47058824],
 [0.38039216],
 [0.22352941],
 ...,
 [0.23529412],
 [0.53725493],
 [0.70588237]],

 [0.37254903],
 [0.29019609],
 [0.14509805],
 ...,
 [0.19607843],
 [0.50980395],
 [0.68235296]],

 ...,

 [[0.37254903],
   [0.29019609],
   [0.14509805],
   ...,
   [0.19607843],
   [0.50980395],
   [0.68235296]],

 [0.47058824],
 [0.38039216],
 [0.22352941],
 ...,
 [0.23529412],

```

```
[0.53725493],
[0.70588237]],

[[0.79215688],
 [0.63921571],
 [0.36470589],
 ...,
 [0.57647061],
 [0.67843139],
 [0.73725492]]],

[[[0.79215688],
  [0.63921571],
  [0.36470589],
  ...,
  [0.57647061],
  [0.67843139],
  [0.73725492]]],

[[0.47058824],
 [0.38039216],
 [0.22352941],
 ...,
 [0.23529412],
 [0.53725493],
 [0.70588237]]],

[[0.37254903],
 [0.29019609],
 [0.14509805],
 ...,
 [0.19607843],
 [0.50980395],
 [0.68235296]]],

...,

[[0.37254903],
 [0.29019609],
 [0.14509805],
 ...,
 [0.19607843],
 [0.50980395],
 [0.68235296]]],

[[0.47058824],
 [0.38039216],
 [0.22352941],
 ...,
 [0.23529412],
 [0.53725493],
 [0.70588237]]],

[[0.79215688],
 [0.63921571],
 [0.36470589],
 ...,
 [0.57647061],
 [0.67843139],
 [0.73725492]]],
```

```

[[[0.79215688],
  [0.63921571],
  [0.36470589],
  ...,
  [0.57647061],
  [0.67843139],
  [0.73725492]]],

[[[0.47058824],
  [0.38039216],
  [0.22352941],
  ...,
  [0.23529412],
  [0.53725493],
  [0.70588237]]],

[[[0.37254903],
  [0.29019609],
  [0.14509805],
  ...,
  [0.19607843],
  [0.50980395],
  [0.68235296]]],

...,

[[[0.37254903],
  [0.29019609],
  [0.14509805],
  ...,
  [0.19607843],
  [0.50980395],
  [0.68235296]]],

[[[0.47058824],
  [0.38039216],
  [0.22352941],
  ...,
  [0.23529412],
  [0.53725493],
  [0.70588237]]],

[[[0.79215688],
  [0.63921571],
  [0.36470589],
  ...,
  [0.57647061],
  [0.67843139],
  [0.73725492]]]]])}

```


Predicted: X350AB20

True: X350AB20



Input img




Activations



Predicted: Y664PK26
True: Y664PK26



Predicted: M427AE55
True: M427AE55




Predicted: E201XE22
True: E201XE22



Input img



Activations



Predicted: C618PA58
True: C618PA58



Predicted: X568YK58
True: X568YK58

Input img




Predicted: E440CX78
True: E440CX78

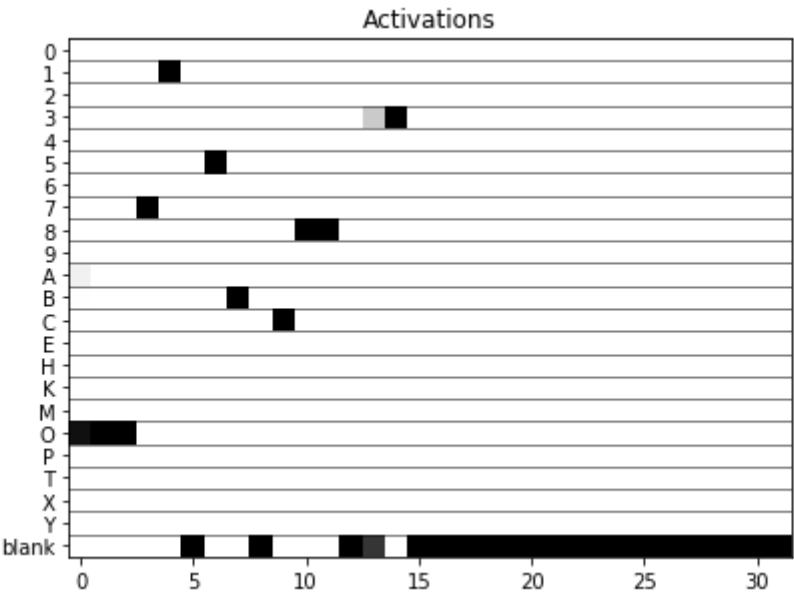

Input img



Activations



Predicted: 0715BC83
True: 0715BC83





In [12]:

```
def adjust_gamma(image, gamma=1.0):  
    # build a lookup table mapping the pixel values [0, 255] to  
    # their adjusted gamma values  
    invGamma = 1.0 / gamma  
    table = np.array([((i / 255.0) ** invGamma) * 255  
                      for i in np.arange(0, 256)]).astype("uint8")  
    # apply gamma correction using the lookup table  
    return cv2.LUT(image, table)
```



In [13]:

```
model.summary()  
model.save("/data/mplate2_model.h5")
```

Layer (type) cted to	Output Shape	Param #	Conne
=====			
the_input (InputLayer)	(None, 128, 64, 1)	0	
=====			
conv1 (Conv2D) nput[0][0]	(None, 128, 64, 16)	160	the_i
=====			
max1 (MaxPooling2D) [0][0]	(None, 64, 32, 16)	0	conv1
=====			
conv2 (Conv2D) [0][0]	(None, 64, 32, 16)	2320	max1
=====			
max2 (MaxPooling2D) [0][0]	(None, 32, 16, 16)	0	conv2
=====			
reshape (Reshape) [0][0]	(None, 32, 256)	0	max2
=====			
dense1 (Dense) pe[0][0]	(None, 32, 32)	8224	resha
=====			
gru1 (GRU) 1[0][0]	(None, 32, 512)	837120	dense
=====			
gru1_b (GRU) 1[0][0]	(None, 32, 512)	837120	dense
=====			
add_1 (Add) [0][0]	(None, 32, 512)	0	gru1 gru1_ b[0][0]
=====			
gru2 (GRU) [0][0]	(None, 32, 512)	1574400	add_1
=====			
gru2_b (GRU)	(None, 32, 512)	1574400	add_1

[0][0]

concatenate_1 (Concatenate) [0][0]	(None, 32, 1024)	0	gru2 gru2_ b[0][0]
dense2 (Dense) tenate_1[0][0]	(None, 32, 23)	23575	conca
softmax (Activation) 2[0][0]	(None, 32, 23)	0	dense
the_labels (InputLayer)	(None, 8)	0	
input_length (InputLayer)	(None, 1)	0	
label_length (InputLayer)	(None, 1)	0	
ctc (Lambda) ax[0][0]	(None, 1)	0	softm the_l input label
abels[0][0]			
_length[0][0]			
_length[0][0]			
=====			
=====			
Total params: 4,857,319			
Trainable params: 4,857,319			
Non-trainable params: 0			



In [14]:


```
def decode_single_image(img, model):
    img_w = 128
    img_h = 64
    img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    img = cv2.resize(img, (img_w, img_h))
    #plt.imshow(img)
    img = img.astype(np.float32)
    img /= 255
    net_inp = model.get_layer(name='the_input').input
    net_out = model.get_layer(name='softmax').output
    img = img.T
    img = np.expand_dims(img, -1)
    X_data = np.ones([1, img_w, img_h, 1])
    X_data[0] = img
    net_out_value = sess.run(net_out, feed_dict={net_inp:X_data})
    return decode_batch(net_out_value)
```



In [15]:

```
# decode some png file
img = cv2.imread("/data/real_numbers/fine.png")
plt.imshow(img)
display(decode_single_image(img, model))
```

['A674YH33']





In [16]:

```
def rotate_bound(image, angle):  
    # grab the dimensions of the image and then determine the  
    # center  
    (h, w) = image.shape[:2]  
    (cX, cY) = (w // 2, h // 2)  
  
    # grab the rotation matrix (applying the negative of the  
    # angle to rotate clockwise), then grab the sine and cosine  
    # (i.e., the rotation components of the matrix)  
    M = cv2.getRotationMatrix2D((cX, cY), -angle, 1.0)  
    cos = np.abs(M[0, 0])  
    sin = np.abs(M[0, 1])  
  
    # compute the new bounding dimensions of the image  
    nW = int((h * sin) + (w * cos))  
    nH = int((h * cos) + (w * sin))  
  
    # adjust the rotation matrix to take into account translation  
    M[0, 2] += (nW / 2) - cX  
    M[1, 2] += (nH / 2) - cY  
  
    # perform the actual rotation and return the image  
    return cv2.warpAffine(image, M, (nW, nH))
```



In [17]:


```
import PIL
img = PIL.Image.open("/data/real_numbers/1.jpg")
display(img)
```



In [18]:

```
# work on image above
img = cv2.imread("/data/real_numbers/1n.png")
img = adjust_gamma(img, 8.0)
img = rotate_bound(img, 7)
plt.imshow(img)
img = cv2.resize(img, (128, 64))
img = img[14:49, 0:128]
plt.imshow(img)
display(decode_single_image(img, model))
```

['C']






In [23]:

```
img = cv2.imread("/data/real_numbers/2.jpg")
plt.imshow(img)
```

Out[23]:

<matplotlib.image.AxesImage at 0x7f16907bbba8>



In [58]:

```
img = cv2.imread("/data/real_numbers/2.jpg")
#img = adjust_gamma(img, 2.0)
img = rotate_bound(img, 1)
#img = img[1390:1900, 900:3300]
img = img[1395:1395+128*4, 970:970+64*36]

img = cv2.resize(img, (128, 64))
plt.imshow(img)
display(decode_single_image(img, model))
```

['K3']





In [60]:

```
img = cv2.imread("/data/real_numbers/fine.png")  
img = cv2.resize(img, (128, 64))  
plt.imshow(img)  
display(decode_single_image(img, model))
```

['A674YH33']

